

INTRO TO NEURAL NETWORKS

Elizabeth Heon

5/19/2025




INTRO

What, where, why?

What is a Neural Network?

A neural network is a machine learning model formed of interconnected layers of nodes, called neurons, which are designed to mimic the function of biological neurons

- Neurons are connected together in layers, forming networks. The output of one neuron becomes input to another.
- The connections in networks are associated with **weights**, which determine how much the output of one neuron influences
- Networks are **trained** by adjusting the weights to improve performance



Where are Neural Networks Used?

4


- Image recognition
- Speech recognition
- Recommendation Algorithm

NETFLIX

- **Many applications in the sciences:**
 - Predict protein structure from amino acids
 - Predict chemical compounds with useful properties
 - Find atoms in EM datasets
 - Etc.

AlphaFold

Accelerating breakthroughs in biology with AI



M. Ziatdinov, ACS Nano 2017

Neural networks are used in many applications in the physical sciences, technology, etc.



Types of Learning:

5

- Supervised Learning: Labelled training data is available – network learns to predict the correct answer for new data
 - Regression
 - Classification
 - Semantic and Instance Segmentation
- Unsupervised Learning: No labelled training data provided – networks learn intrinsic patterns in the data
 - Clustering
- Semi-supervised Learning: Some labelled training data is provided


This talk will focus on supervised learning, where a network is supplied labelled training data



Regression and Classification

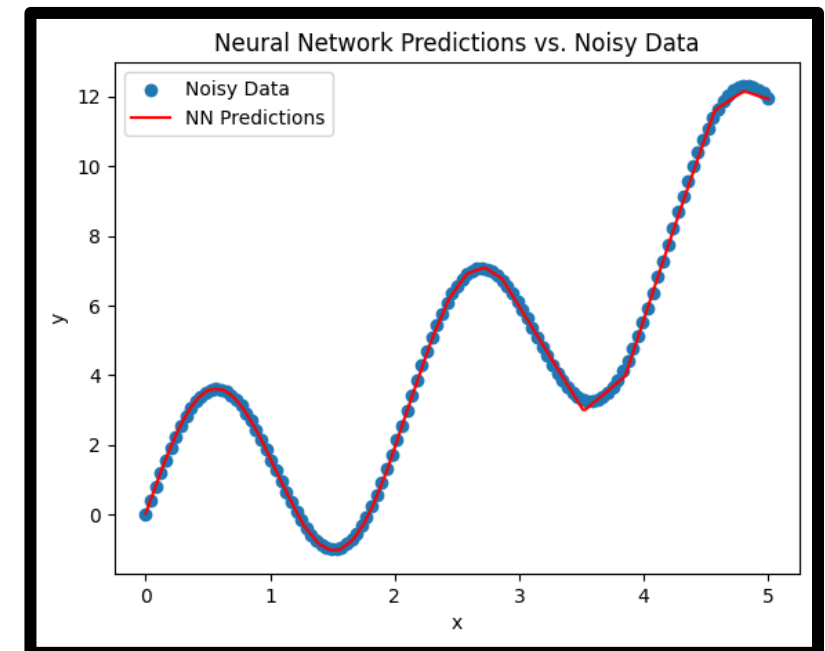
Classification

- Classification tasks -> Output of network is categorical
- Examples:
 - Medical Diagnosis: Input -> medical imaging, Output -> Diagnosis
 - Character Recognition: Input -> Image of handwritten character, Output -> Letter A-Z or digit 0-9



Regression


- Regression tasks -> Output of network is numerical
- Examples:
 - Linear regression / function interpolation



I. NEURONS

Biological Neurons:

8



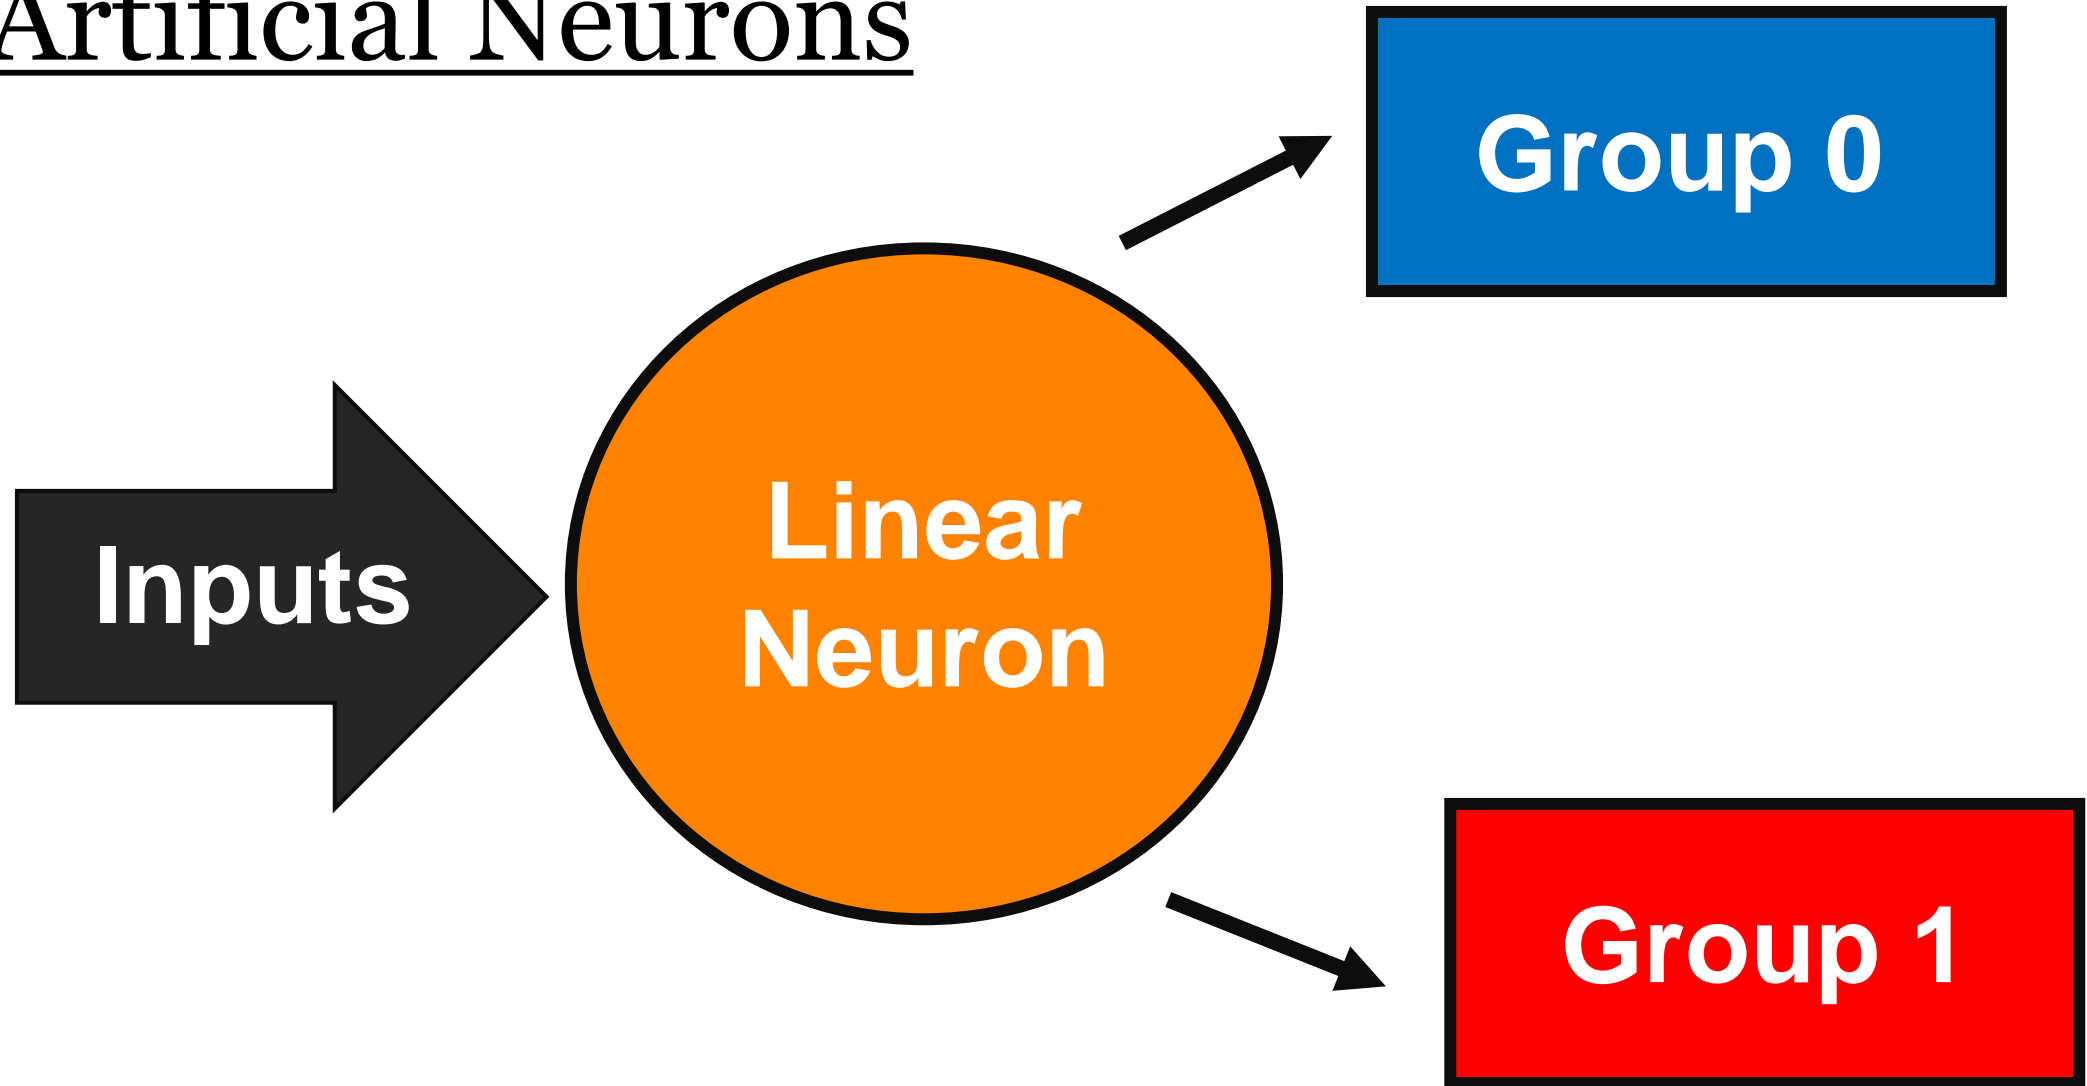
- Dendrites receive messages from other neurons – these messages can be excitatory (encourage the neuron to fire) or inhibitory (prevent firing)
- Depending on the combination of inputs, the neuron may fire an “action potential” which travels down the axon to other neurons

A (biological) neuron takes a combination of inputs and decides which output to produce – potential or no potential



Artificial Neurons

9



Consider a binary linear neuron – given a set of inputs, the neuron can output either 0 or 1 -> essentially, the neuron can classify a set of inputs into one of two groups



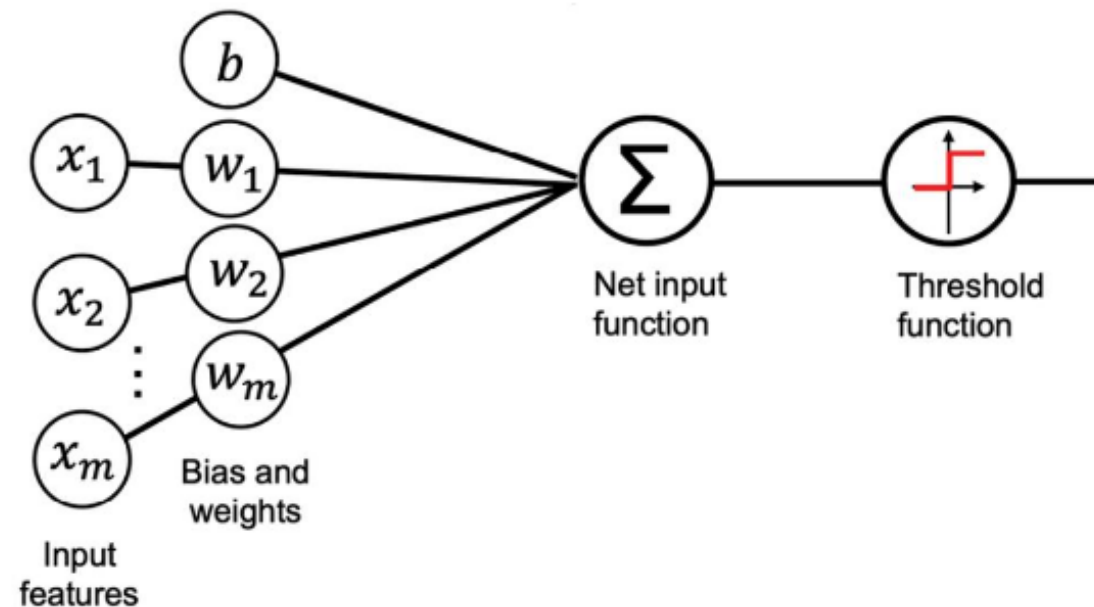
Artificial Neurons

- Neuron receives a vector of **inputs**, x_i
- Inputs are multiplied by vector of **weights** to form net input function:

$$z = \sum_i x_i w_i$$

- The net input is compared to a value called the **threshold** (θ):

- Output:** $\sigma(z) = \begin{cases} 1 & \text{if } z > \theta \\ 0 & \text{otherwise} \end{cases}$



Artificial neurons take in a set of inputs and determine what output to produce using a set of weights and a threshold



Bias

11

- The critical point between inputs that result in output 0 and inputs that result in output 1 is given by:

$$z = \sum_i x_i w_i = \theta$$

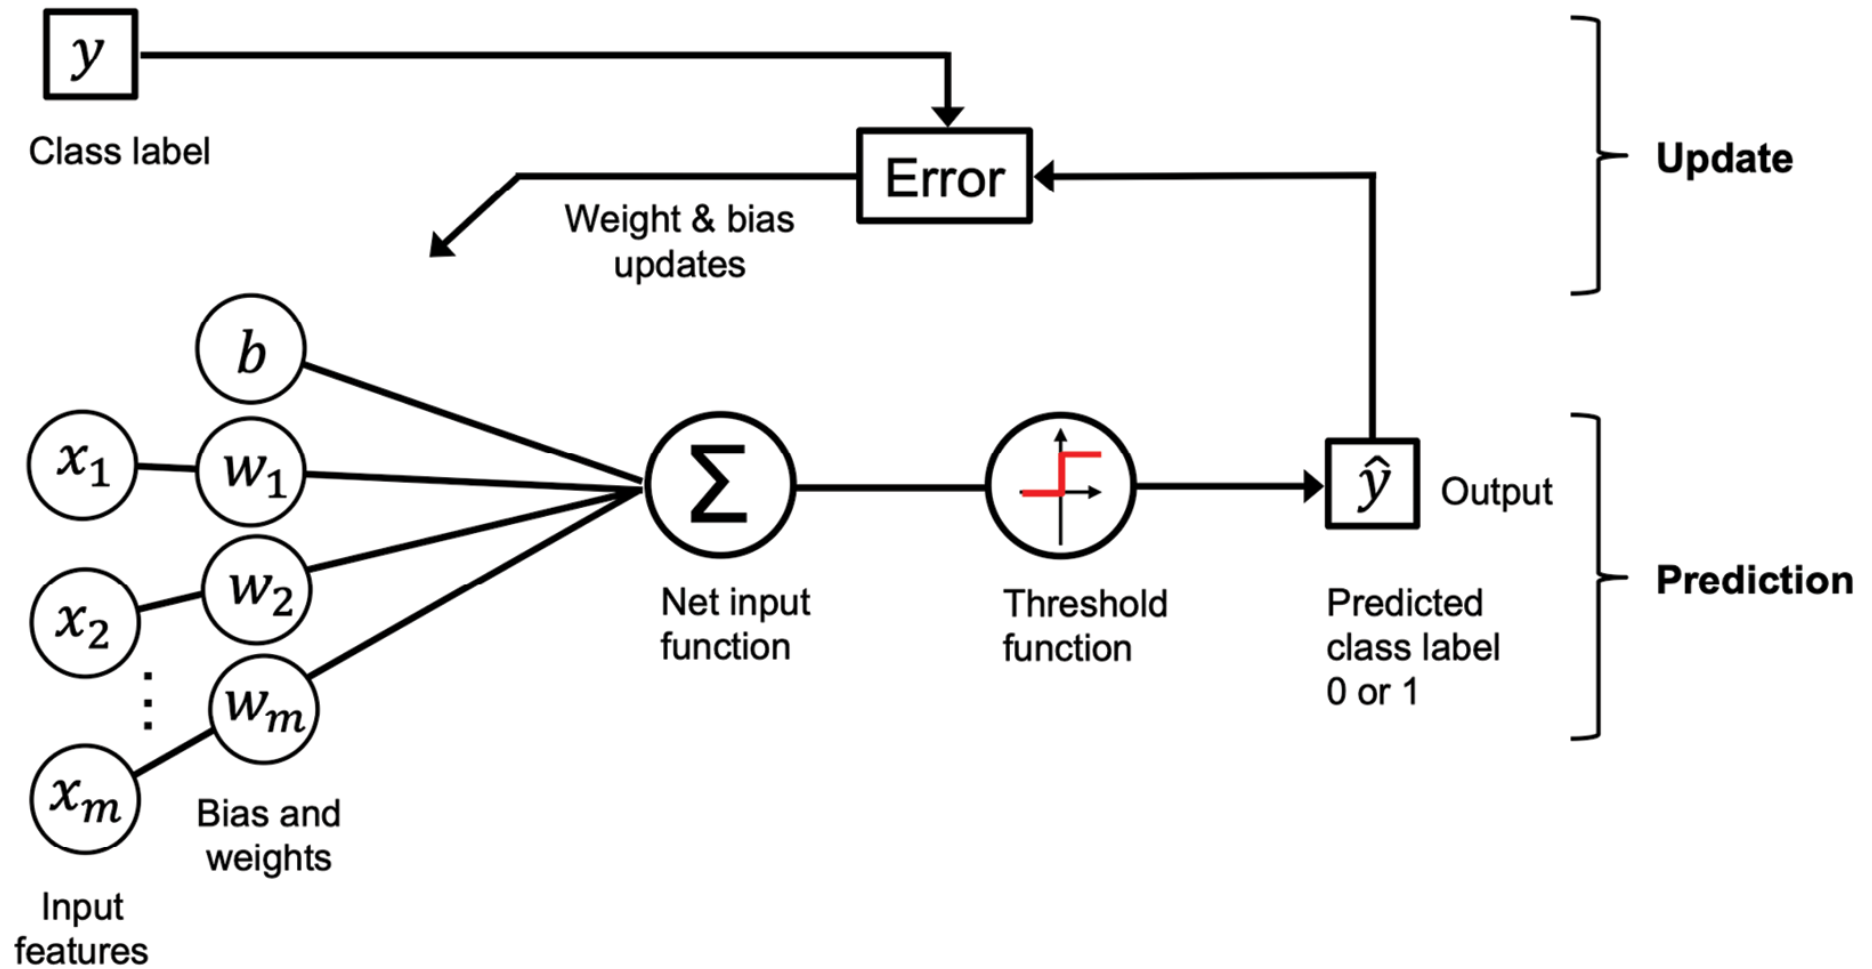
- Subtract the value of θ from both sides:

$$\begin{aligned} x_i w_i + \dots + x_n w_n - \theta &= 0 \\ x_i w_i + \dots + x_n w_n + b &= 0 \end{aligned}$$

- The term b is known as the **bias** – using the bias makes the threshold value trainable along with the weights



Training a Linear Neuron - Overview



Training is the process of updating the weights and bias to improve the performance of the model. Training is performed iteratively



Training Algorithm - Verbal

13

- To train the network we require a set of training data. Training data consists of:
 - Vectors of inputs $\rightarrow x_i$
 - The “**ground truth**” (i.e. correct) output for each set of inputs, $y^{(i)}$
- For each vector in the training, g set, calculate the **model prediction**, $\hat{y}^{(i)}$, then:
 - If the prediction is correct, do nothing
 - If the prediction is wrong, update the weights and bias to move the decision boundary in a way that reduces error
 - Repeat until the bias and weights stop changing



Training Algorithm - Math

14

1) Initialize the bias and weights either to zero or to small random numbers



2) For each example in the training set:

- **Compute the output using the current bias and weights:** $\hat{y}^{(i)} = w^T x^{(i)} + b$
- **Update the weights and the bias:**
 - $w_j = w_j + \Delta w$ where $\Delta w = \eta(y^{(i)} - \hat{y}^{(i)})x_j^{(i)}$
 - $b = b + \Delta b$ where $\Delta b = \eta(y^{(i)} - \hat{y}^{(i)})$



3) Repeat until prediction matches output for all examples in training set



The Problem of Separability


- Consider a linear neuron that takes two inputs, x_1 and x_2 :
- We thus have: $x_1 w_1 + x_2 w_2 = \theta$
- Rearrange:

$$x_2 w_2 = \theta - x_1 w_1$$

$$x_2 = \frac{\theta}{w_2} - \frac{w_1}{w_2} x_1$$

$$x_2 = \mathbf{a}x_1 + \mathbf{b}$$

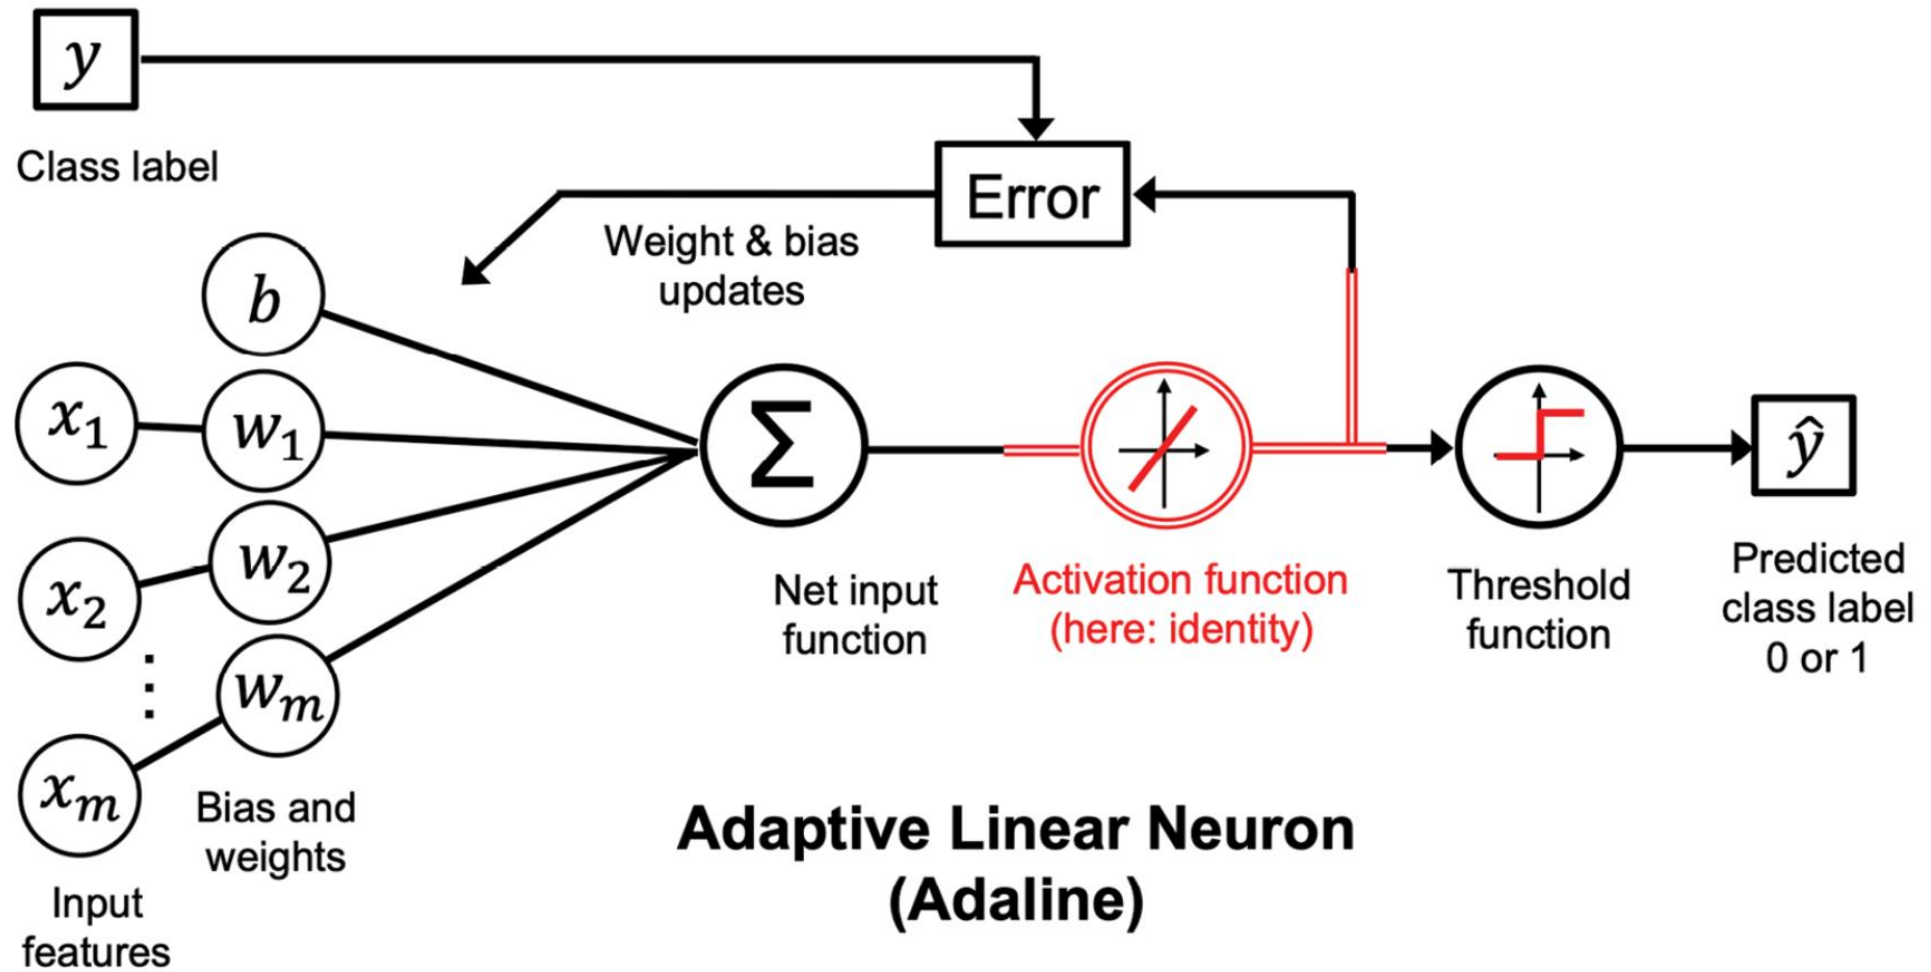
- Results in linear equation – **decision line/plane**



Single neurons can only solve problems that are linearly separable







Adaptive Linear Neuron (Adaline)



Activation Functions:

- Activation functions introduce **non-linearity** to neural networks – otherwise even deep networks can only solve linear problems
- Non-linearity \rightarrow output has a non-linear relationship to input – decision surface can be curved




Activation functions introduce non-linearity to neural networks, allowing networks to learn complex patterns in data




Choosing Activation Functions

ReLU Activation Function



- **Rectified Linear Unit:** If input is negative, output is zero. If input is positive, output = input
- Often used for image data or in other cases where negative inputs are not physically meaningful

Leaky ReLU Activation Function




- **Leaky ReLU:** Introduce a small non-zero slope for negative inputs
- Helps avoid “dying ReLU” where negative inputs cause a node in a network to become inactive




Choosing Activation Functions

Sigmoid Activation Function



- **Sigmoid:** Maps output to values between 0 and 1
- Ideal for situations where output should be interpreted as a probability

Tanh Activation Function



- **Tanh:** Maps output to values between -1 and 1
- Often used in hidden layers



Training an Adaline – Loss Function

- We train an Adaline by minimizing a **loss function**:

$$L(w, b) = \frac{1}{n} \sum_{i=1}^n \left(y^{(i)} - \sigma(z^{(i)}) \right)^2$$

Model Prediction (arrow pointing to $\sigma(z^{(i)})$)
Ground Truth (arrow pointing to $y^{(i)}$)

- A loss function quantifies the difference between the network prediction and the ground truth
- The above is called the Mean Squared Error (MSE) loss and is common for regression problems


Loss functions quantify error for Adaline; training occurs by minimizing the loss



Gradient Descent - Analogy

21

- For every combination of weights and bias there is a corresponding value of the loss. This forms a loss surface/hypersurface. To train the network we must move to the minimum of the surface



- You are on a hill, and want to reach the valley (minimum) nearby
- To reach the minimum, make a series of steps, always making sure to **move down hill**



Gradient Descent - Math

22

- To move “downhill” we need the slope of the loss function
 - To find the slope of a function, take the derivative (in this case **partial derivatives** since loss is a multivariate function)

1) Start at an arbitrary position on the loss surface.


2) Use the learning rule:

$$\Delta w = -\eta \nabla_w L(w, b)$$

$$\Delta b = -\eta \nabla_b L(w, b)$$

$$w := w + \Delta w$$

$$b := w + \Delta b$$




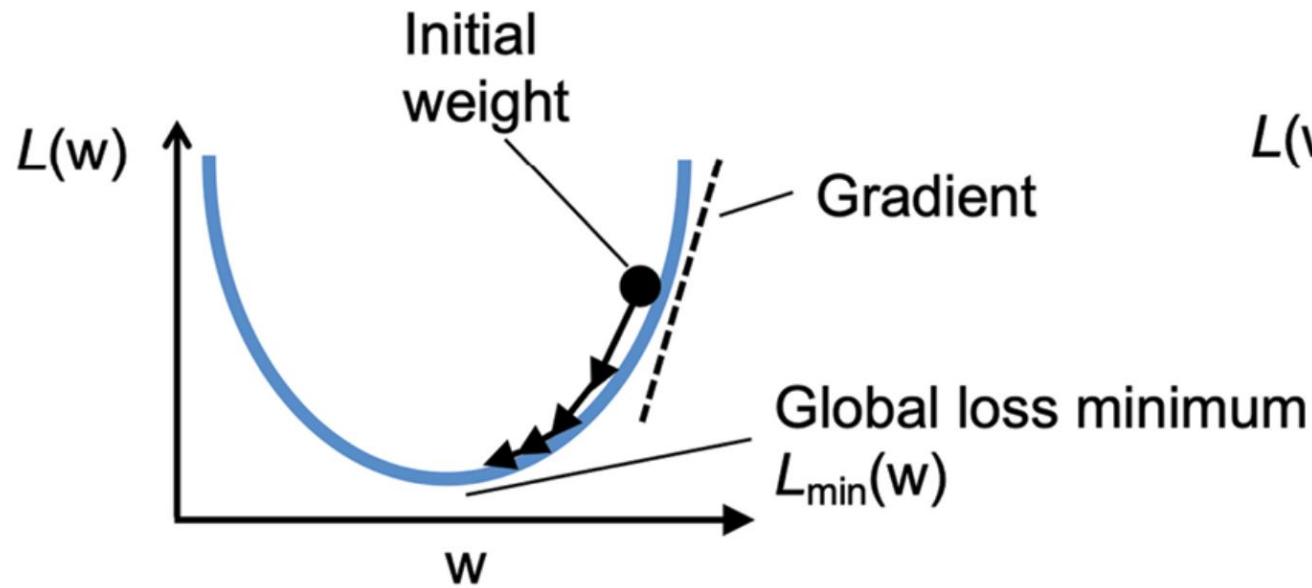
η = Learning Rate

Determines how large each “step” is



Role of Learning Rate


- Learning rate, η , determines how quickly weights and bias change – how large each “step” is. Too small and the network will take a very long time to reach the minimum. Too large and it may overshoot the minimum and fail to converge.



II. FROM NEURONS TO NETWORK

Putting Neurons Together

- Deep architectures allow networks to learn complex patterns in data (so long as they have non-linear activation functions) – how does a multi-layer network work?



- Composed of multiple layers of artificial neurons – the inner layers are known as “**hidden layers**”
 - The number of nodes per layer and the connectivity between nodes varies between networks
- Each node receives inputs, applies weights and biases, and passes outputs onto the next layer



Training a Multilayer Network

26

- To train a multi-layer network, each neurons weights and bias must be updated. To train multi-layer networks an approach known as backpropagation is used.
- Backpropagation consists of a forward and backward pass:
 - **Forward pass:** put training data into network at input layer, and propagate forward through the network, producing an output.
 - Calculate the loss based on the network prediction and the ground truth
 - **Backward pass:** Compute the gradient of the loss function with respect to each neuron's weights by chain rule
 - Update weights

Multi-layer networks are trained by backpropagation, which moves layer by layer to determine appropriate updates to the network weights



Once a network is trained, the performance can be evaluated according to several possible metrics, including:

- **Accuracy:** Defined as the ratio of correctly predicted observations to the total observations.
- **Precision:** The ratio of correctly predicted positive observations to the total predicted positive observations. Important in scenarios where the cost of a false positive is high.
- **Recall (Sensitivity):** The ratio of correctly predicted positive observations to all observations in the actual class. Crucial in situations where missing a positive instance is costly.



III. USING NETWORKS

There are Many Types of Neural Networks

- Multilayer perceptron
- Convolutional Neural Networks
- Autoencoders
- Physics informed neural networks
- Etc.

Given the many types of networks that are possible, how do you know which type to use?



Where to Start:

30

Problem-Centric Approach:

- Always start with the problem at hand: Analyze the nature of input data and desired output
- Choose a network architecture that aligns with the type and structure of your data
- Select a loss function that reflects the objective of the problem
- Metrics should be chosen based on what measures success for your specific task

Hyperparameter Tuning:

- Once the architecture and loss function are set, proceed to tune hyperparameters including network structure, optimizers, regularization, etc.
- Hyperparameter tuning should be guided by the chosen metrics and the nature of the problem



Coding Task



Thank you